

## Keras -- MLPs on MNIST

In [2]:

```
# if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use this command
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
```

In [3]:

```
%matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

In [4]:

```
# the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

In [5]:

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)  
Number of training examples : 10000 and each image is of shape (28, 28)

In [6]:

```
# if you observe the input shape its 3 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

In [7]:

```
# after converting the input images from 3d to 2d vectors

print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d)"%(X_train.shape[1]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d)"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)  
Number of training examples : 10000 and each image is of shape (784)

In [8]:

```
# An example data point
```

```
# An example data point
print(X_train[0])
```

[	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	3	18	18	18	126	136	175	26	166	255	
247	127	0	0	0	0	0	0	0	0	0	0	0	0	0	30	36	94	154	
170	253	253	253	253	225	172	253	242	195	64	0	0	0	0	0	0	0	0	
	0	0	0	0	0	49	238	253	253	253	253	253	253	253	251	93	82		
82	56	39	0	0	0	0	0	0	0	0	0	0	0	0	0	18	219	253	
253	253	253	253	198	182	247	241	0	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	80	156	107	253	253	205	11	0	43	154		
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0	14	1	154	253	90	0	0	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	0	0	0	139	253	190	2	0	
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	11	190	253	70	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	35	241	
225	160	108	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	81	240	253	253	119	25	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0	0	45	186	253	253	150	27	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	0	0	0	16	93	252	253	187	
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	249	253	249	64	0	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	0	0	0	0	46	130	183	253	
253	207	2	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	
	0	0	0	0	39	148	229	253	253	253	250	182	0	0	0	0	0	0	
	0	0	0	0	0	0	0	0	0	0	0	0	24	114	221	253	253	253	
253	201	78																	

```
# if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the data
#  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$ 

X_train = X_train/255
X_test = X_test/255
```

```
X_train = X_train/255
X_test = X_test/255
```

```
# example data point after normlizing
print(X_train[0])
```

[illegible]

[illegible]

[illegible]

In [11]:

```
# here we are having a class number for each image
print("Class label of first image :", y_train[0])

# lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ",Y_train[0])
```

```
Class label of first image : 5
After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]
```

## Softmax classifier

In [12]:

```
# https://keras.io/getting-started/sequential-model-guide/

# The Sequential model is a linear stack of layers.
# you can create a Sequential model by passing a list of layer instances to the constructor:

# model = Sequential([
#     Dense(32, input_shape=(784,)),
#     Activation('relu'),
#     Dense(10),
```

```

#     Activation('softmax'),
# ])

# You can also simply add layers via the .add() method:

# model = Sequential()
# model.add(Dense(32, input_dim=784))
# model.add(Activation('relu'))

###

# https://keras.io/layers/core/

# keras.layers.Dense(units, activation=None, use_bias=True, kernel_initializer='glorot_uniform',
# bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None,
# activity_regularizer=None,
# kernel_constraint=None, bias_constraint=None)

# Dense implements the operation: output = activation(dot(input, kernel) + bias) where
# activation is the element-wise activation function passed as the activation argument,
# kernel is a weights matrix created by the layer, and
# bias is a bias vector created by the layer (only applicable if use_bias is True).

# output = activation(dot(input, kernel) + bias) => y = activation(WT. X + b)

####

# https://keras.io/activations/

# Activations can either be used through an Activation layer, or through the activation argument s
# upported by all forward layers:

# from keras.layers import Activation, Dense

# model.add(Dense(64))
# model.add(Activation('tanh'))

# This is equivalent to:
# model.add(Dense(64, activation='tanh'))

# there are many activation functions ar available ex: tanh, relu, softmax

from keras.models import Sequential
from keras.layers import Dense, Activation

```

In [13]:

```

# some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20

```

## MLP + ReLU + ADAM with 2 Layers

In [14]:

```

# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras
from keras.layers.normalization import BatchNormalization
from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))

```

```
seed=None))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 512)	401920
batch_normalization_1 (Batch Normalization)	(None, 512)	2048
dropout_1 (Dropout)	(None, 512)	0
dense_2 (Dense)	(None, 128)	65664
batch_normalization_2 (Batch Normalization)	(None, 128)	512
dropout_2 (Dropout)	(None, 128)	0
dense_3 (Dense)	(None, 10)	1290
Total params: 471,434		
Trainable params: 470,154		
Non-trainable params: 1,280		

In [15]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 8s 137us/step - loss: 0.4862 - acc: 0.8509 -
val_loss: 0.1624 - val_acc: 0.9497
Epoch 2/20
60000/60000 [=====] - 8s 141us/step - loss: 0.2521 - acc: 0.9245 -
val_loss: 0.1230 - val_acc: 0.9632
Epoch 3/20
60000/60000 [=====] - 7s 109us/step - loss: 0.2041 - acc: 0.9388 -
val_loss: 0.1000 - val_acc: 0.9696
Epoch 4/20
60000/60000 [=====] - 8s 135us/step - loss: 0.1719 - acc: 0.9479 -
val_loss: 0.0915 - val_acc: 0.9712
Epoch 5/20
60000/60000 [=====] - 9s 153us/step - loss: 0.1554 - acc: 0.9534 -
val_loss: 0.0873 - val_acc: 0.9709
Epoch 6/20
60000/60000 [=====] - 7s 114us/step - loss: 0.1381 - acc: 0.9583 -
val_loss: 0.0824 - val_acc: 0.9734
Epoch 7/20
60000/60000 [=====] - 6s 106us/step - loss: 0.1308 - acc: 0.9608 -
val_loss: 0.0772 - val_acc: 0.9760
Epoch 8/20
60000/60000 [=====] - 7s 109us/step - loss: 0.1188 - acc: 0.9637 -
val_loss: 0.0720 - val_acc: 0.9768
Epoch 9/20
60000/60000 [=====] - 7s 109us/step - loss: 0.1142 - acc: 0.9658 -
val_loss: 0.0717 - val_acc: 0.9794
Epoch 10/20
60000/60000 [=====] - 10s 168us/step - loss: 0.1047 - acc: 0.9671 - val_l
oss: 0.0678 - val_acc: 0.9790
Epoch 11/20
60000/60000 [=====] - 8s 139us/step - loss: 0.1010 - acc: 0.9693 -
val_loss: 0.0680 - val_acc: 0.9787
Epoch 12/20
60000/60000 [=====] - 8s 129us/step - loss: 0.0976 - acc: 0.9698 -
val_loss: 0.0639 - val_acc: 0.9808
```

```

Epoch 13/20
60000/60000 [=====] - 7s 113us/step - loss: 0.0931 - acc: 0.9714 -
val_loss: 0.0627 - val_acc: 0.9801
Epoch 14/20
60000/60000 [=====] - 8s 129us/step - loss: 0.0834 - acc: 0.9747 -
val_loss: 0.0610 - val_acc: 0.9811
Epoch 15/20
60000/60000 [=====] - 8s 125us/step - loss: 0.0844 - acc: 0.9743 -
val_loss: 0.0605 - val_acc: 0.9823
Epoch 16/20
60000/60000 [=====] - 8s 137us/step - loss: 0.0767 - acc: 0.9754 -
val_loss: 0.0603 - val_acc: 0.9814
Epoch 17/20
60000/60000 [=====] - 8s 138us/step - loss: 0.0786 - acc: 0.9751 -
val_loss: 0.0591 - val_acc: 0.9833
Epoch 18/20
60000/60000 [=====] - 7s 116us/step - loss: 0.0737 - acc: 0.9767 -
val_loss: 0.0596 - val_acc: 0.9816
Epoch 19/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0708 - acc: 0.9782 -
val_loss: 0.0585 - val_acc: 0.9830
Epoch 20/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0700 - acc: 0.9779 -
val_loss: 0.0652 - val_acc: 0.9821

```

In [16]:

```

score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
print("Large Error: %.2f%%" % (100-score[1]*100))

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidaion_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.histrory we will have a list of length equal to number of epochs


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)


```

```

Test score: 0.06517606317649478
Test accuracy: 0.9821
Large Error: 1.79%

```





In [17]:


```
w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



## MLP + ReLU + ADAM with 3 Layers



In [18]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(360, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 512)	401920
batch_normalization_3 (Batch Normalization)	(None, 512)	2048
dropout_3 (Dropout)	(None, 512)	0
dense_5 (Dense)	(None, 360)	184680
batch_normalization_4 (Batch Normalization)	(None, 360)	1440
dropout_4 (Dropout)	(None, 360)	0
dense_6 (Dense)	(None, 128)	46208
batch_normalization_5 (Batch Normalization)	(None, 128)	512
dropout_5 (Dropout)	(None, 128)	0
dense_7 (Dense)	(None, 10)	1290
Total params: 638,098		
Trainable params: 636,098		
Non-trainable params: 2,000		

In [19]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 11s 190us/step - loss: 0.7516 - acc: 0.7659 - val_loss: 0.2080 - val_acc: 0.9365
Epoch 2/20
60000/60000 [=====] - 10s 165us/step - loss: 0.3401 - acc: 0.8989 - val_loss: 0.1570 - val_acc: 0.9515
Epoch 3/20
60000/60000 [=====] - 10s 166us/step - loss: 0.2640 - acc: 0.9219 - val_loss: 0.1278 - val_acc: 0.9612
Epoch 4/20
```

```

60000/60000 [=====] - 10s 169us/step - loss: 0.2219 - acc: 0.9353 - val_l
oss: 0.1143 - val_acc: 0.9667
Epoch 5/20
60000/60000 [=====] - 10s 174us/step - loss: 0.1948 - acc: 0.9426 - val_l
oss: 0.1046 - val_acc: 0.9672
Epoch 6/20
60000/60000 [=====] - 11s 176us/step - loss: 0.1777 - acc: 0.9482 - val_l
oss: 0.0908 - val_acc: 0.9729
Epoch 7/20
60000/60000 [=====] - 11s 192us/step - loss: 0.1576 - acc: 0.9527 - val_l
oss: 0.0874 - val_acc: 0.9745
Epoch 8/20
60000/60000 [=====] - 10s 164us/step - loss: 0.1490 - acc: 0.9553 - val_l
oss: 0.0828 - val_acc: 0.9764
Epoch 9/20
60000/60000 [=====] - 12s 197us/step - loss: 0.1359 - acc: 0.9593 - val_l
oss: 0.0798 - val_acc: 0.9773
Epoch 10/20
60000/60000 [=====] - 15s 252us/step - loss: 0.1289 - acc: 0.9612 - val_l
oss: 0.0787 - val_acc: 0.9770
Epoch 11/20
60000/60000 [=====] - 12s 208us/step - loss: 0.1211 - acc: 0.9642 - val_l
oss: 0.0754 - val_acc: 0.9776
Epoch 12/20
60000/60000 [=====] - 12s 198us/step - loss: 0.1177 - acc: 0.9646 - val_l
oss: 0.0707 - val_acc: 0.9788
Epoch 13/20
60000/60000 [=====] - 13s 217us/step - loss: 0.1106 - acc: 0.9672 - val_l
oss: 0.0730 - val_acc: 0.9779
Epoch 14/20
60000/60000 [=====] - 13s 224us/step - loss: 0.1057 - acc: 0.9680 - val_l
oss: 0.0695 - val_acc: 0.9796
Epoch 15/20
60000/60000 [=====] - 13s 223us/step - loss: 0.0999 - acc: 0.9698 - val_l
oss: 0.0713 - val_acc: 0.9796
Epoch 16/20
60000/60000 [=====] - 14s 227us/step - loss: 0.0940 - acc: 0.9714 - val_l
oss: 0.0651 - val_acc: 0.9816
Epoch 17/20
60000/60000 [=====] - 15s 251us/step - loss: 0.0900 - acc: 0.9725 - val_l
oss: 0.0670 - val_acc: 0.9801
Epoch 18/20
60000/60000 [=====] - 13s 223us/step - loss: 0.0862 - acc: 0.9737 - val_l
oss: 0.0653 - val_acc: 0.9816
Epoch 19/20
60000/60000 [=====] - 13s 218us/step - loss: 0.0831 - acc: 0.9749 - val_l
oss: 0.0637 - val_acc: 0.9817
Epoch 20/20
60000/60000 [=====] - 11s 178us/step - loss: 0.0820 - acc: 0.9746 - val_l
oss: 0.0648 - val_acc: 0.9815

```

In [20]:

```

score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
print(" Error: %.2f%%" % (100-score[1]*100))

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va
lidation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

```

```
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.06477750879065133  
 Test accuracy: 0.9815  
 Error: 1.85%



In [21]:

```
w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

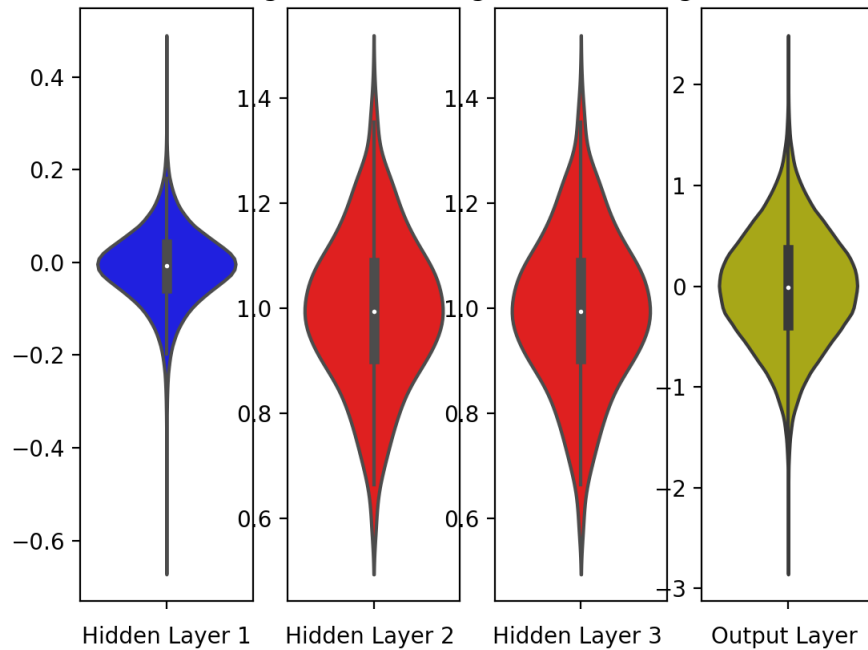
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

Trained model Weights model Weights model Weights model Weights



## MLP + ReLU + ADAM with 5 Layers

In [22]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(620, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(480, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(350, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(220, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Layer (type)	Output Shape	Param #
dense_8 (Dense)	(None, 620)	486700
batch_normalization_6 (Batch Normalization)	(None, 620)	2480

batch_normalization_6 (Batch Normalization)	(None, 480)	1920
dropout_6 (Dropout)	(None, 620)	0
dense_9 (Dense)	(None, 480)	298080
batch_normalization_7 (Batch Normalization)	(None, 480)	1920
dropout_7 (Dropout)	(None, 480)	0
dense_10 (Dense)	(None, 350)	168350
batch_normalization_8 (Batch Normalization)	(None, 350)	1400
dropout_8 (Dropout)	(None, 350)	0
dense_11 (Dense)	(None, 220)	77220
batch_normalization_9 (Batch Normalization)	(None, 220)	880
dropout_9 (Dropout)	(None, 220)	0
dense_12 (Dense)	(None, 128)	28288
batch_normalization_10 (Batch Normalization)	(None, 128)	512
dropout_10 (Dropout)	(None, 128)	0
dense_13 (Dense)	(None, 10)	1290
=====		
Total params: 1,067,120		
Trainable params: 1,063,524		
Non-trainable params: 3,596		

In [23]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 22s 365us/step - loss: 0.9467 - acc: 0.7095 - val_loss: 0.2148 - val_acc: 0.9358
Epoch 2/20
60000/60000 [=====] - 18s 298us/step - loss: 0.3351 - acc: 0.9064 - val_loss: 0.1470 - val_acc: 0.9572
Epoch 3/20
60000/60000 [=====] - 22s 375us/step - loss: 0.2519 - acc: 0.9316 - val_loss: 0.1209 - val_acc: 0.9648
Epoch 4/20
60000/60000 [=====] - 24s 392us/step - loss: 0.2065 - acc: 0.9437 - val_loss: 0.1061 - val_acc: 0.9696
Epoch 5/20
60000/60000 [=====] - 20s 339us/step - loss: 0.1819 - acc: 0.9508 - val_loss: 0.1079 - val_acc: 0.9715
Epoch 6/20
60000/60000 [=====] - 21s 347us/step - loss: 0.1664 - acc: 0.9550 - val_loss: 0.0953 - val_acc: 0.9747
Epoch 7/20
60000/60000 [=====] - 20s 334us/step - loss: 0.1548 - acc: 0.9578 - val_loss: 0.0864 - val_acc: 0.9768
Epoch 8/20
60000/60000 [=====] - 20s 340us/step - loss: 0.1418 - acc: 0.9612 - val_loss: 0.0824 - val_acc: 0.9782
Epoch 9/20
60000/60000 [=====] - 22s 372us/step - loss: 0.1332 - acc: 0.9639 - val_loss: 0.0868 - val_acc: 0.9765
Epoch 10/20
60000/60000 [=====] - 18s 292us/step - loss: 0.1226 - acc: 0.9665 - val_loss: 0.0788 - val_acc: 0.9794
Epoch 11/20
60000/60000 [=====] - 19s 318us/step - loss: 0.1133 - acc: 0.9687 - val_loss: 0.0815 - val_acc: 0.9788
```

```

Epoch 12/20
60000/60000 [=====] - 22s 360us/step - loss: 0.1120 - acc: 0.9690 - val_loss: 0.0739 - val_acc: 0.9810
Epoch 13/20
60000/60000 [=====] - 21s 344us/step - loss: 0.1075 - acc: 0.9701 - val_loss: 0.0760 - val_acc: 0.9809
Epoch 14/20
60000/60000 [=====] - 22s 364us/step - loss: 0.1073 - acc: 0.9715 - val_loss: 0.0746 - val_acc: 0.9802
Epoch 15/20
60000/60000 [=====] - 23s 388us/step - loss: 0.1010 - acc: 0.9721 - val_loss: 0.0743 - val_acc: 0.9796
Epoch 16/20
60000/60000 [=====] - 20s 325us/step - loss: 0.0983 - acc: 0.9735 - val_loss: 0.0698 - val_acc: 0.9832
Epoch 17/20
60000/60000 [=====] - 19s 323us/step - loss: 0.0901 - acc: 0.9754 - val_loss: 0.0671 - val_acc: 0.9824
Epoch 18/20
60000/60000 [=====] - 19s 315us/step - loss: 0.0842 - acc: 0.9768 - val_loss: 0.0666 - val_acc: 0.9829
Epoch 19/20
60000/60000 [=====] - 20s 333us/step - loss: 0.0844 - acc: 0.9758 - val_loss: 0.0689 - val_acc: 0.9824
Epoch 20/20
60000/60000 [=====] - 20s 340us/step - loss: 0.0882 - acc: 0.9764 - val_loss: 0.0626 - val_acc: 0.9838

```

In [24]:

```

score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
print(" Error: %.2f%%" % (100-score[1]*100))

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)


```

Test score: 0.06262384752128274

Test accuracy: 0.9838

Error: 1.62%





In [25]:

```
w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
h5_w = w_after[8].flatten().reshape(-1,1)
out_w = w_after[10].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 6, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 6, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')


plt.subplot(1, 6, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')


plt.subplot(1, 6, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='r')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 6, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h5_w, color='r')
plt.xlabel('Hidden Layer 5 ')

plt.subplot(1, 6, 6)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```

Trained model Weights





## Results

In [30]:

```
from prettytable import PrettyTable
x = PrettyTable()

names = ["Relu with 2 layers", "Relu with 3 layers", "Relu with 5 layers"]
Test_Accuracy = [97.91, 97.61, 97.8]
Test_Score = [0.0691, 0.080, 0.082]
numbering = [1, 2, 3]
ptable = PrettyTable()
# Adding columns
ptable.add_column("S.NO.", numbering)
ptable.add_column("MODEL", names)

ptable.add_column("Test Score", Test_Score)
ptable.add_column("Test Accuracy", Test_Accuracy)
# Printing the Table
print(ptable)
```

S.NO.	MODEL	Test Score	Test Accuracy
1	Relu with 2 layers	0.0691	97.91
2	Relu with 3 layers	0.08	97.61
3	Relu with 5 layers	0.082	97.8

## Reference

In [ ]:

<https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/exercise-try-different-mlpa-architectures-on-mnist-dataset/>